

# Priority Queues (Heaps)

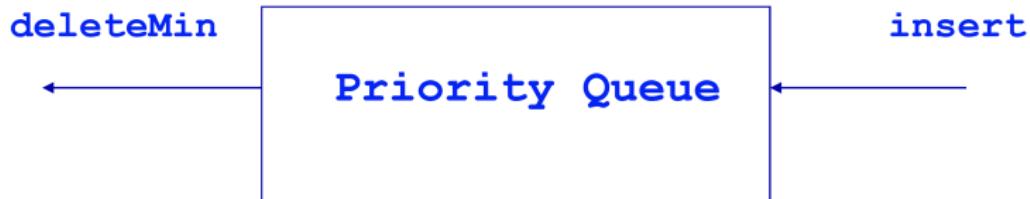
November 2, 2023

# Priority Queues

- Many applications require that we process records with keys **in order**, but not necessarily in full sorted order.
- Often we collect a set of items and process the one with the current minimum value. e.g.
  - Jobs sent to a printer,
  - Operating system job scheduler in a multi-user environment.
  - Simulation environments (Discrete Event Simulator)
- An appropriate data structure for this purpose is called a **priority queue**.

# Definition

- A priority queue is a data structure that supports two basic operations:
  - insert a new item and
  - remove the minimum item.



# Implementations Options

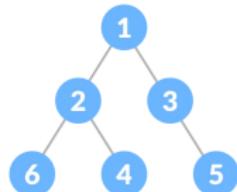
- A simple linked list:
  - Insertion at the front  $O(1)$ ; `deleteMin`  $O(N)$ , or
  - Keep list sorted; insertion  $O(N)$ , `deleteMin`  $O(1)$
- A binary search tree (BST):
  - This gives an  $\Theta(\log N)$  average for both operations.
  - But BST class supports many operations that are not required.
- Binary Heap:
  - Can be implemented as a simple array and does not require links.
  - Supports both operations in  $O(\log N)$  worst-case time.
- d-Heaps:
  - A more general case in which a parent node can have  $d$  children
- Leftist and Skew Heaps
- Binomial Queues

# Binary Heap

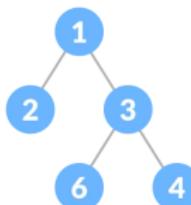
- The binary heap is the **classic** method used to implement priority queues.
- We use the term **heap** to refer to the binary heap.
- Heap is different from the term heap used in dynamic memory allocation.
- Heap has two properties:
  - Structure property
  - Ordering property

# Structure Property

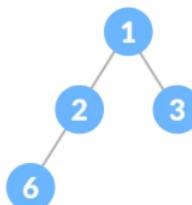
- A **heap** is a complete binary tree.
- A **complete binary tree** is a binary tree that is completely filled, with the possible exception of the bottom level, which is filled from left to right.



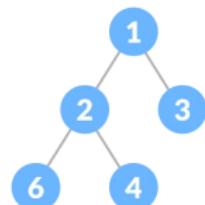
✗ Full Binary Tree  
✗ Complete Binary Tree



✓ Full Binary Tree  
✗ Complete Binary Tree



✗ Full Binary Tree  
✓ Complete Binary Tree

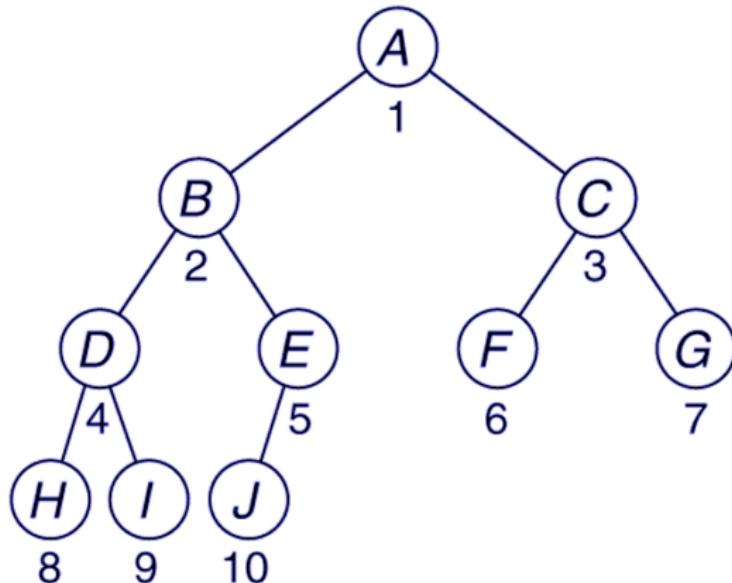


✓ Full Binary Tree  
✓ Complete Binary Tree

# Properties of a complete binary tree

- A complete binary tree of height  $h$  has between  $2^h$  and  $2^{h+1} - 1$  nodes
- The height of a complete binary tree is  $\lfloor \log N \rfloor$ .
- It can be implemented as an array such that:
  - For any element in array position  $i$  :
    - the left child is in position  $2i$ ,
    - the right child is in the cell after the left child ( $2i + 1$ ), and
    - the parent is in position  $\lfloor i/2 \rfloor$ .

# A complete binary tree and its array representation



# Heap-Order Property

- In a heap, for every node  $X$  with parent  $P$ , the key in  $P$  is **smaller than or equal to** the key in  $X$ .
- Thus the minimum element is always at the root.
  - Thus we get the operation `findMin` in constant time.
- A **max heap** supports access of the maximum element instead of the minimum, by changing the heap property slightly.

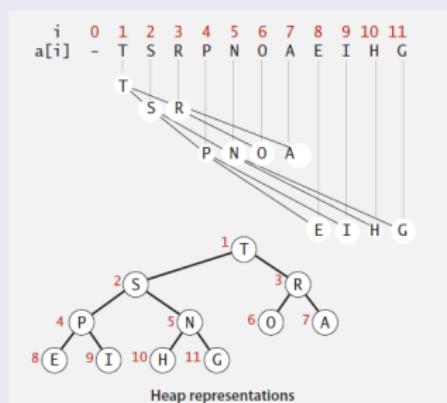
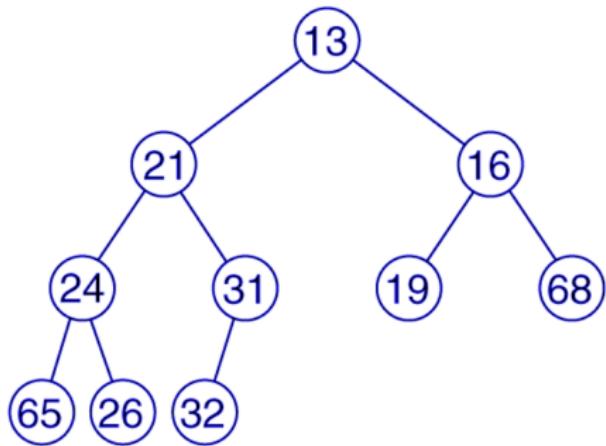
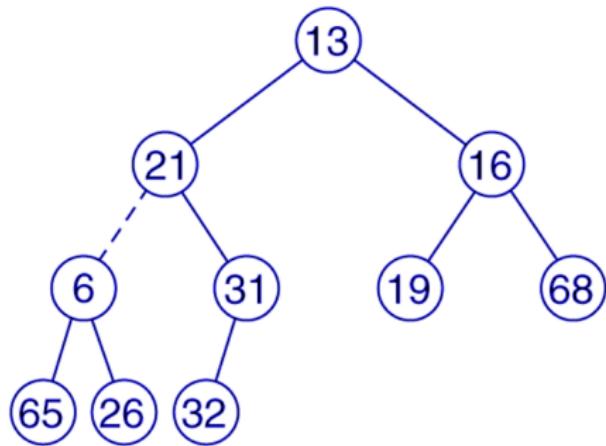


Figure: Max heap example

## Two complete trees



(a)



(b)

(a) a heap;

(b) not a heap

# Binary Heap Class

```
public class BinaryHeap<AnyType extends Comparable<? super AnyType>>
{
    public BinaryHeap( ){...}
    public BinaryHeap( int capacity ){...}
    public BinaryHeap( AnyType [ ] items ){...}

    public void insert( AnyType x ){...}
    public AnyType findMin( ){...}
    public AnyType deleteMin( ){...}
    public boolean isEmpty( ){...}
    public void makeEmpty( ){...}

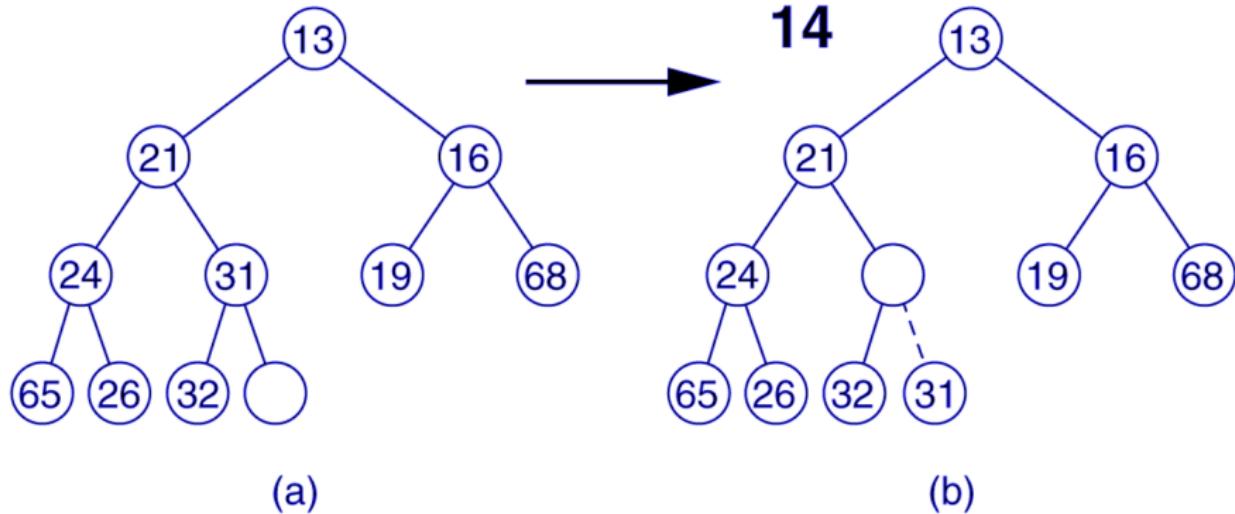
    private static final int DEFAULT_CAPACITY = 10;
    private int currentSize; // Number of elements in heap
    private AnyType [ ] array; // The heap array

    private void percolateDown( int hole ){...}
    private void buildHeap( ){...}
    private void enlargeArray( int newSize ){...}
}
```

# Basic Heap Operations: Insert

- To insert an element  $X$  into the heap:
  - We create a hole in the next available location.
  - If  $X$  can be placed there without violating the heap property, then we do so and are done.
  - Otherwise
    - We bubble up the hole toward the root by sliding the element in the hole's parent down.
    - We continue this until  $X$  can be placed in the hole.
- This general strategy is known as a **percolate up**.

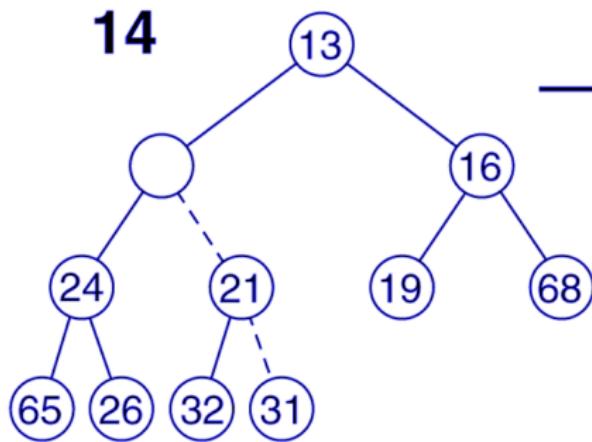
Attempt to insert 14, creating the hole and bubbling the hole up



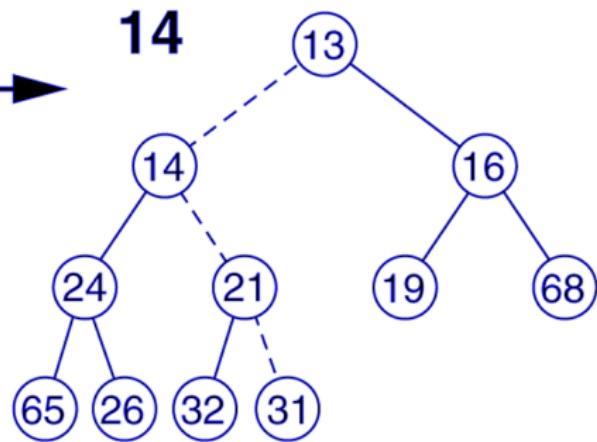
(a)

(b)

The remaining two steps required to insert 14



(a)



(b)

# Insert procedure

```
/**  
 * Insert item x, allowing duplicates.  
 */  
public void insert( AnyType x )  
{  
    if( currentSize == array.length - 1 )  
        enlargeArray( array.length * 2 + 1 );  
  
    // Percolate up  
    int hole = ++currentSize;  
  
    for(array[ 0 ] = x; x.compareTo(array[hole/2]) < 0; hole/=2)  
        array[ hole ] = array[ hole / 2 ];  
    array[ hole ] = x;  
}
```

Why do we assign x to the position 0?

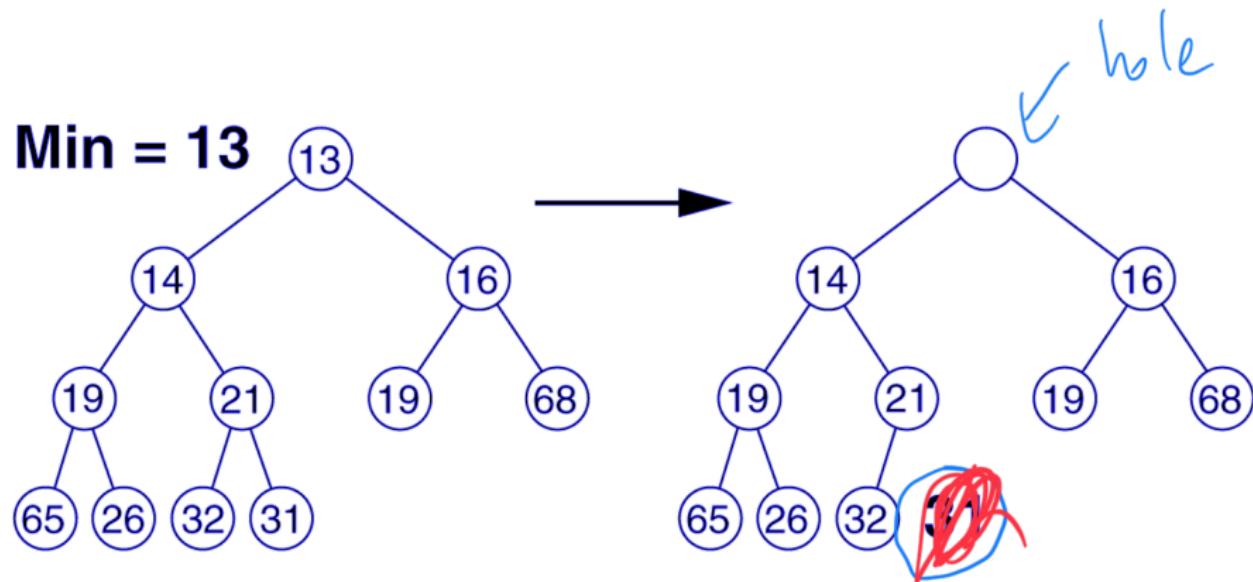
## Analysis of Insert

- The time to do insertion could be as much as  $O(\log N)$  if the inserted element is the new minimum.
- On average, the percolation up terminates early and average runtime for insert is constant.

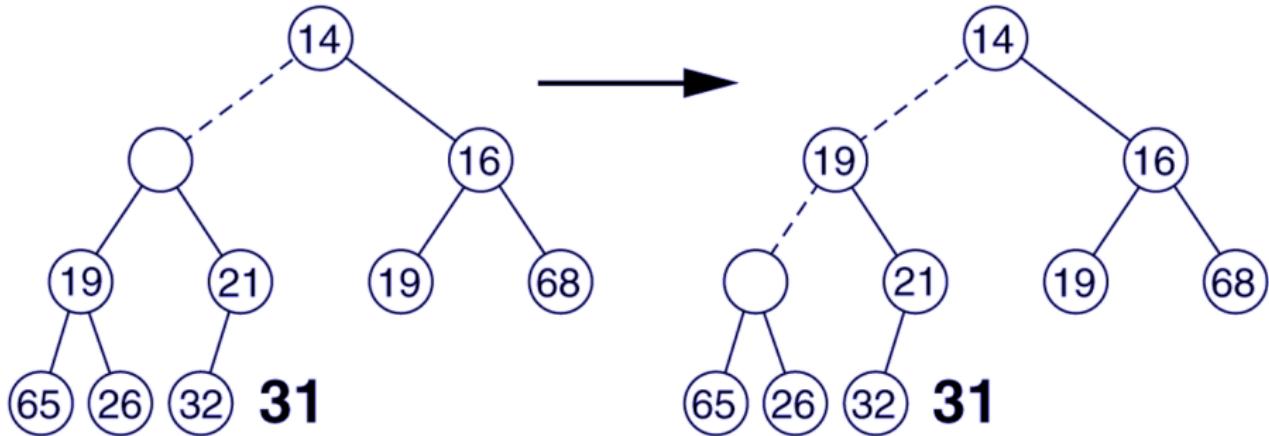
## Delete Minimum

- `deleteMin` is handled in a similar manner as insertion:
- Remove the minimum; a hole is created at the root.
- The last element X must move to somewhere in the heap.
  - If X can be placed in the hole then we are done.
  - Otherwise,
    - We slide the smaller of the hole's children into the hole, thus pushing the hole one level down.
    - We repeat this until X can be placed in the hole.

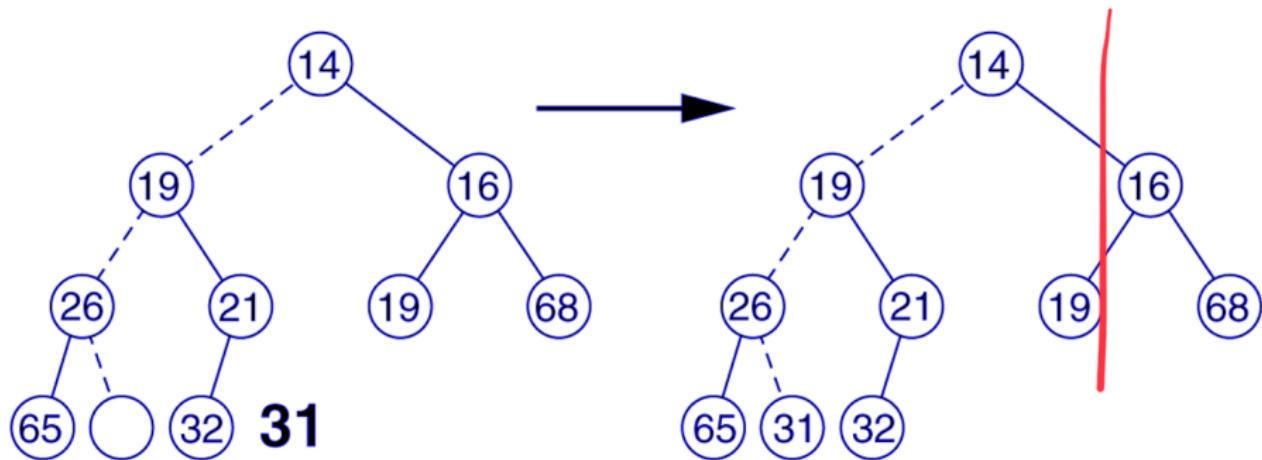
## Creation of the hole at the root



## The next two steps in the deleteMin operation



## The last two steps in the deleteMin operation



## deleteMin procedure

```
// Remove the smallest item from the priority queue.  
// Throw UnderflowException if empty.  
public AnyType deleteMin( )  
{  
    if( isEmpty( ) )  
        throw new UnderflowException( );  
  
    AnyType minItem = findMin( ); // Array location 1  
  
    array[ 1 ] = array[ currentSize-- ];  
  
    percolateDown( 1 );  
  
    return minItem;  
}
```

## percolateDown procedure

```
/***
 * Internal method to percolate down in the heap.
 * hole is the index at which the percolate begins.
 */
private void percolateDown( int hole )
{
    int child;
    AnyType tmp = array[ hole ];

    for( ; hole * 2 <= currentSize; hole = child )
    {
        child = hole * 2;
        if(child != currentSize && array[child+1].compareTo(array[child]) < 0)
            child++;
        if( array[child].compareTo( tmp ) < 0 )
            array[hole] = array[child];
        else
            break;
    }
    array[ hole ] = tmp;
}
```

# Analysis of Delete Minimum

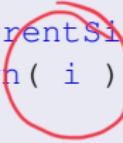
- The worst case running time for `deleteMin` is  $O(\log N)$
- On average, the element placed at the root is percolated down almost to the bottom of the heap (its original level), so the average running time is  $O(\log N)$

# Building a Heap

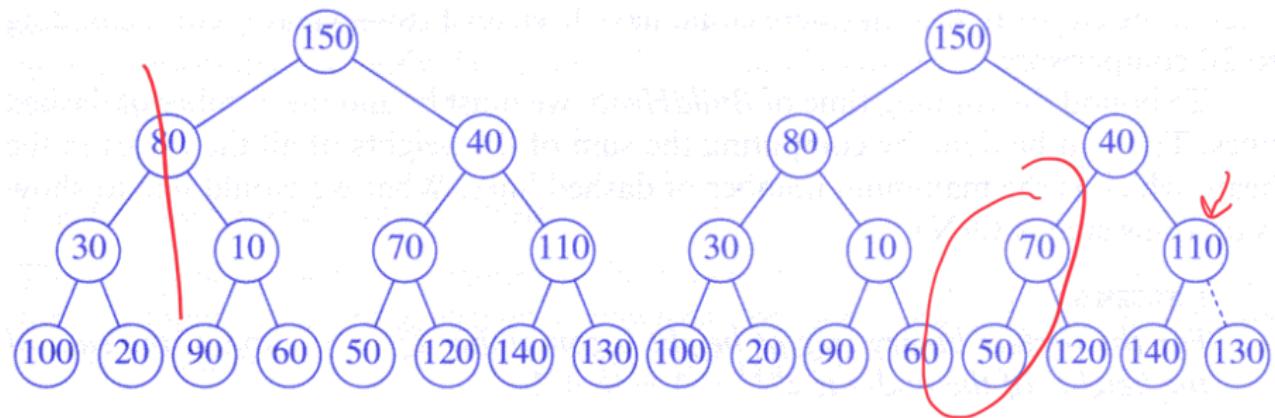
- Take as input  $N$  items and place them into an empty heap.
- Obviously this can be done with  $N$  successive inserts:  $O(N \log N)$  worst case,  $O(N)$  on average.
- However `buildHeap` operation can be done in linear time ( $O(N)$ ) by applying a percolate down routine to nodes in reverse level order.

## buildHeap method

```
/**  
 * Establish heap order property from an arbitrary  
 * arrangement of items. Runs in linear time.  
 */  
private void buildHeap( )  
{  
    for( int i = currentSize / 2; i > 0; i-- )  
        percolateDown( i );  
}
```



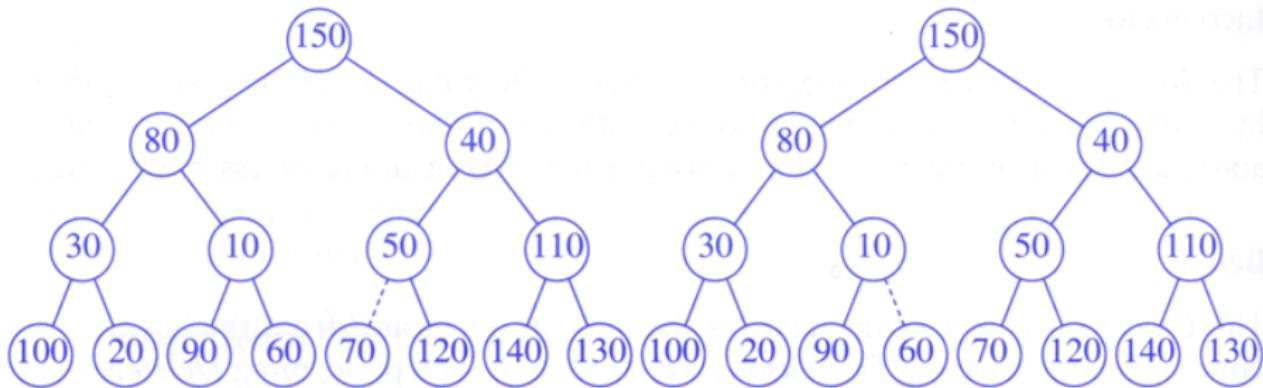
# Implementation of the linear-time BuildHeap method



Initial heap

After percolateDown(7)

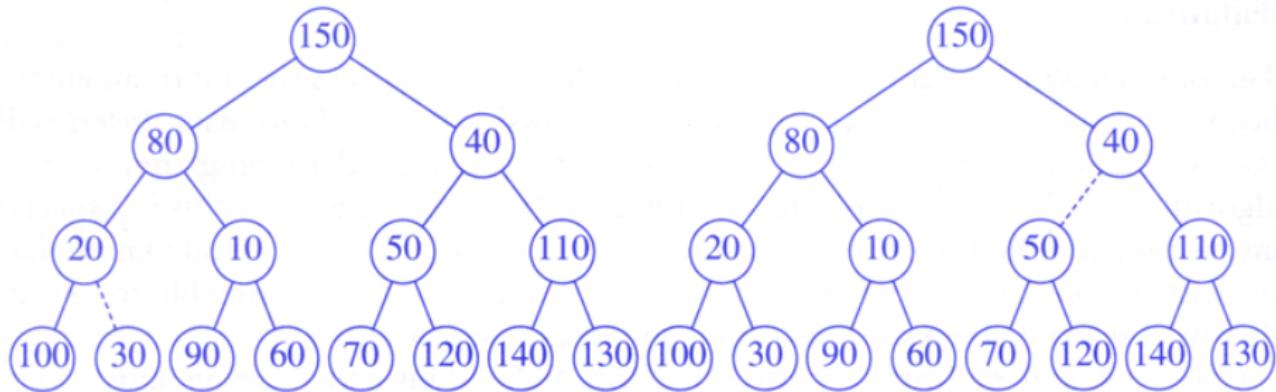
## BuildHeap (Cont.)



(a) After `percolateDown(6);`

(b) after `percolateDown(5)`

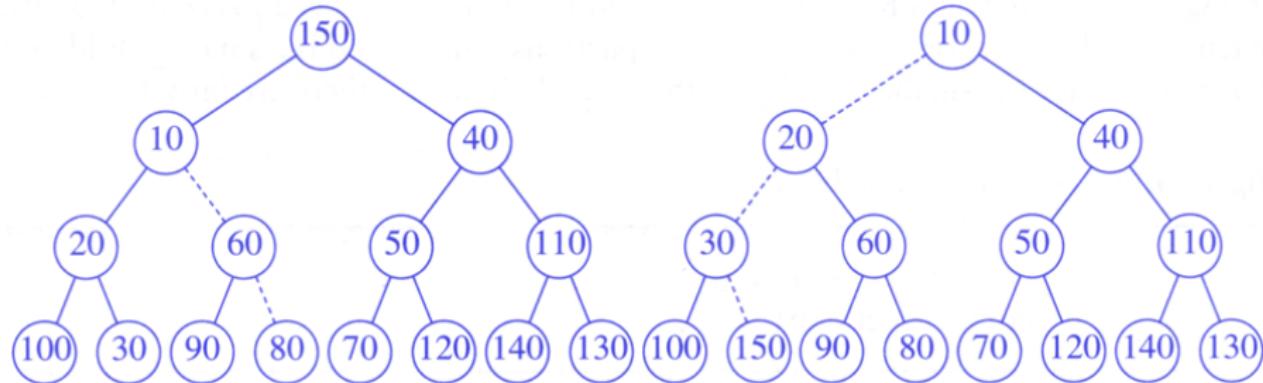
## BuildHeap (Cont.)



(a) After `percolateDown(4);`

(b) after `percolateDown(3)`

## BuildHeap (Cont.)



(a) After `percolateDown(2)`; (b) after `percolateDown(1)` and  
BuildHeap terminates

# Analysis of BuildHeap

- The linear time bound of BuildHeap, can be shown by computing the sum of the heights of all the nodes in the heap
- For the perfect binary tree of height  $H$  containing  $N = 2^{H+1} - 1$  nodes, the sum of the heights of the nodes is  $N - H - 1$ .
- Thus it is  $O(N)$ .
- 1 node at height  $H$
- 2 nodes at height  $H - 1$
- $2^2$  nodes at height  $H - 2 \dots$
- $S = \sum_{i=0}^H 2^i (H - i)$

# Some Applications of Priority Queues

- Operating system design
- Implementation of several graph algorithms efficiently
- The selection problem (and sorting)
- Discrete Event Simulation

# Selection Problem

- Given a list of  $N$  elements, and an integer  $k$ , the selection problem is to find the  $k$ th smallest element.
  - Read  $N$  elements into an array, apply *buildHeap* algorithm to this array. Then perform  $k$  *deleteMin* operations. The last element extracted from the heap is the answer.
  - The runtime is  $O(N + k \log N)$ .
  - Note that if we run the algorithm for  $k = N$  and record the values as they leave the heap, we will have sorted list that will run in  $O(N \log N)$  time. This is known as *heapsort*.
  - We can find the  $k$ th largest element using max heap.

# Discrete Event Simulation

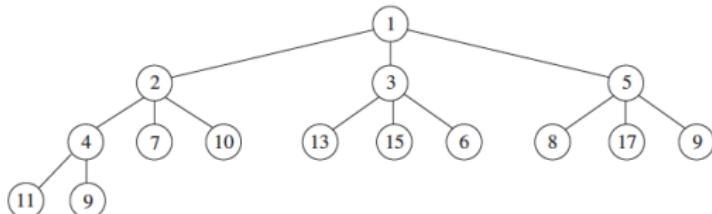
- For some queuing problems analytical analysis is complex hence simulation might be needed.
- In our simulations, we can use a quantum unit which we will refer as *tick*. One option is to start simulation clock at zero ticks and advance the clock one tick at a time to see if there is an event.
  - The problem is that runtime does not depend on customers or events but number of ticks, which is not part of the input.
- Instead, we can advance the clock to the next event time at each stage. To find the events that are nearest in the feature, it is appropriate to keep them in a priority queue.

## DES Example: Emergency Room Simulation

- Patients **arrive** at time  $t$  with injury of criticality  $C$ 
  - If no patients are waiting and there is a free doctor, assign them to doctor and create a future **departure** event; else put patient in the Criticality priority queue
- Patient departs at time  $t$ 
  - If someone in Criticality queue, pull out most critical and assign to doctor; create a future **departure** event

## d-Heaps

- A generalization which is exactly like a binary heap except that all nodes have  $d$  children, hence a binary heap is  $2\text{-heap}$



- The tree will be shallower, hence inserts will be improved compared to binary heap.
- Due to  $d - 1$  comparisons among the children, `deleteMin` will be more expensive.
- Unless  $d$  is a power of 2, multiplications and divisions to find children and parents will be much more expensive (can not be done by bit shift).

# Leftist Heaps and Skew Heaps and Binomial Queues

- Any array based implementation of priority queue will have  $\Theta(N)$  to **merge** (combine two heaps into one) two equal sized heaps.
- For efficient implementation of merge, we need linked data structure.  
Note that this will make all other operations slower.
- Leftist Heaps, Skew Heaps and Binomial Queues are linked based priority queue implementations that support efficient merge operation.

# Leftist Heap

## Null Path Length Definition

We define the null path length,  $npl(X)$ , of any node  $X$  to be the length of the shortest path from  $X$  to a node without two children. Thus, the  $npl$  of a node with zero or one child is 0, while  $npl(null) = -1$ .

## Leftist Heap Property

The leftist heap property is that for every node  $X$  in the heap, the null path length of the left child is at least as large as that of the right child.

## Observation

- Leftist heaps tend to have deep left paths
- In contrast to a binary heap (which is always a complete binary tree), a leftist tree may be very unbalanced.
- The general idea for the leftist heap operations is to perform all the work on the right path, which is guaranteed to be short

# Leftist Heap Example



Figure 6.20 Null path lengths for two trees; only the left tree is leftist

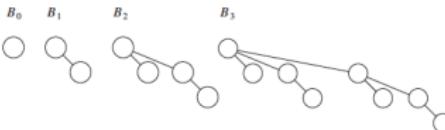
- Heavier on left side :  $npl(\text{right}(i)) \leq npl(\text{left}(i))$ .
- The shortest path to a descendant node is through the right child. Every subtree is also a leftist tree and  $npl(i) = 1 + npl(\text{right}(i))$ .
- The path from root to rightmost leaf is the shortest path from root to a leaf.
- If the path to rightmost leaf has  $x$  nodes, then leftist heap has at least  $2^x - 1$  nodes. This means the length of path to rightmost leaf is  $O(\log n)$  for a leftist heap with  $n$  nodes.

# Leftist Heap Merge Algorithm

Visualization at: <https://www.geeksforgeeks.org/leftist-tree-leftist-heap/>  
Merge can be performed in  $O(\log N)$  time

# Binomial Queues

- A binomial queue is not a heap-ordered tree but a collection of heap-ordered trees, known as **forest**.
- Each of the heap-ordered trees is of a constrained form known as a binomial tree
- There is at most one binomial tree of every height.
- A binomial tree of height 0 is one node tree. Binomial tree,  $B_k$  of height  $k$  is constructed by attaching a binomial tree  $B_{k-1}$  to the root of another binomial tree  $B_{k-1}$ .
- In other words, a Binomial Tree of order  $k$  can be constructed by taking two binomial trees of order  $k-1$  and making one the rightmost child or the other (leftmost)



# Binomial Queue Construction

Binomial trees of height  $k$  have exactly  $2^k$  nodes.

If we impose heap order on the binomial trees and allow at most one binomial tree of any height, we can represent a priority queue of any size by a collection of binomial trees.

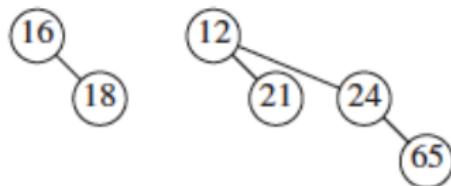
For instance, a priority queue of size 13 could be represented by the forest  $B_3, B_2, B_0$ .

## Advantage:

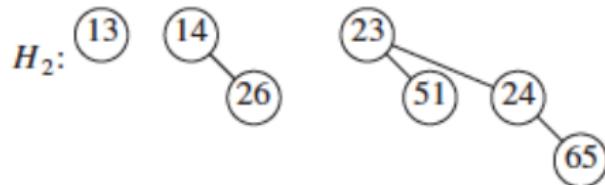
Merging easy with binomial queues in  $O(\log N)$  time compared to  $O(N)$  binary min heap.

# Binomial Queue Example

$H_1:$



$H_2:$



$H_3:$  13

