

Sorting Algorithms

November 15, 2023

Sorting

- **Sorting** is a process that organizes a collection of data into either *ascending* or *descending* order.
- An **internal sort** requires that the collection of data fit entirely in the computer's main memory.
- We can use an **external sort** when the collection of data cannot fit in the computer's main memory all at once but must reside in secondary storage such as on a disk (or tape).
- We will analyze only internal sorting algorithms.

Why Sorting?

- Significant amount of computer output is generally arranged in some **sorted order** so that it can be interpreted.
- Sorting also has indirect uses. An **initial sort** of the data can significantly enhance the performance of an algorithm.
- Majority of programming projects use sorting somewhere, and in many cases, the **sorting cost** determines the running time.
- Applications:
 - Closest pair – Given a set of n numbers, how do you find the pair of numbers that have the smallest difference between them?
 - Frequency distribution – Given a set of n items, which element occurs the largest number of times in the set?
 - Selection – What is the k th largest item in an array?, etc.

Take-Home Lesson

Sorting lies at the heart of many algorithms. Sorting the data is one of the first things any algorithm designer should try in the quest for efficiency.

Preliminaries

- We will discuss the problem of sorting an **array of elements**.
- We assume all array positions contains data to be sorted.
- To simply matters, we will assume in the examples that the array contains **integers**. But any **comparable** type can be used.
- For general sorting algorithms, besides the assignment operator, we will assume **existence of the ">" and "<" operators** which can be used to place a consistent ordering on the input. Sorting under these conditions is known as **"comparison-based sorting"**.
- A comparison-based sorting algorithm makes ordering decisions only on the basis of comparisons.

Sorting Algorithms

- There are many sorting algorithms, such as:
 - Selection Sort
 - Insertion Sort
 - Bubble Sort
 - Shell Sort
 - Heap Sort
 - Merge Sort
 - Quick Sort
- The first three are the foundations for more efficient algorithms.

Insertion Sort

- Insertion sort is a simple sorting algorithm that is appropriate for small inputs.
 - The most common sorting technique used by card players.
- The list is divided into two parts: *sorted* and *unsorted*.
- In each pass, the first element of the unsorted part is picked up, transferred to the sorted sublist, and inserted at the appropriate place.
- A list of n elements will take at most $n - 1$ passes to sort the data.

Insertion Sort Example

Sorted

Unsorted

23	78	45	8	32	56
----	----	----	---	----	----

Original List

23	78	45	8	32	56
----	----	----	---	----	----

After pass 1

23	45	78	8	32	56
----	----	----	---	----	----

After pass 2

8	23	45	78	32	56
---	----	----	----	----	----

After pass 3

8	23	32	45	78	56
---	----	----	----	----	----

After pass 4

8	23	32	45	56	78
---	----	----	----	----	----

After pass 5

Insertion Sort Example

Original	34	8	64	51	32	21	Positions Moved
After $p = 1$	8	34	64	51	32	21	1
After $p = 2$	8	34	64	51	32	21	0
After $p = 3$	8	34	51	64	32	21	1
After $p = 4$	8	32	34	51	64	21	3
After $p = 5$	8	21	32	34	51	64	4

Insertion Sort Algorithm

```
// Simple insertion sort for integer arrays
void insertionSort( int [] a )
{
    for( int p = 1; p < a.length; p++ )
    {
        int tmp = a[p];

        int j;
        for( j = p; j > 0 && tmp < a[j-1]; j-- )
            a[j] = a[j-1];
        a[j] = tmp;
    }
}
```

Insertion Sort – Analysis

- Running time depends on not only the size of the array but also the contents of the array.
- **Best-case:** $\rightarrow O(n)$
 - Array is already sorted in ascending order.
 - Inner loop will not be executed.
 - The number of moves: $2 \times (n - 1) \rightarrow O(n)$
 - The number of key comparisons: $(n - 1) \rightarrow O(n)$
- **Worst-case:** $\rightarrow O(n^2)$
 - Array is in reverse order:
 - Inner loop is executed $i - 1$ times, for $i = 2, 3, \dots, n$
 - The number of moves:
 $2 \times (n - 1) + (1 + 2 + \dots + n - 1) = 2 \times (n - 1) + n \times (n - 1)/2 \rightarrow O(n^2)$
 - The number of key comparisons:
 $(1 + 2 + \dots + n - 1) = n \times (n - 1)/2 \rightarrow O(n^2)$
- **Average-case:** $\rightarrow O(n^2)$
 - We have to look at all possible initial data organizations.
- So, Insertion Sort is $O(n^2)$

Analysis of insertion sort

- Which running time will be used to characterize this algorithm?
 - Best, worst or average?
- Worst:
 - Longest running time (this is the upper limit for the algorithm)
 - It is guaranteed that the algorithm will not be worse than this.
- Sometimes we are interested in the average case. But there are some problems with the average case.
 - It is difficult to figure out the average case. i.e. what is average input?
 - Are we going to assume all possible inputs are equally likely?
 - In fact, **for most sorting algorithms the average case is the same as the worst case.**

Other simple sorting algorithms

Selection Sort

- Scans from left to right
 - In iteration i , find index \min of smallest remaining entry.
 - Swap $a[i]$ and $a[\min]$.
-
- $\text{arr}[] = 64, 25, 12, 22, 11$
 - **1st pass:** 11 25 12 22 64
 - **2nd pass:** 11 12 25 22 64
 - **3rd pass:** 11 12 22 25 64
 - **4th pass:** 11 12 22 25 64
 - **5th pass:** 11 12 22 25 64

Other simple sorting algorithms

Bubble Sort

- repeatedly swap the adjacent elements if they are in the wrong order.

First Pass:

- $(5\ 1\ 4\ 2\ 8) \rightarrow (1\ 5\ 4\ 2\ 8)$, Swaps since $5 > 1$
- $(1\ 5\ 4\ 2\ 8) \rightarrow (1\ 4\ 5\ 2\ 8)$, Swap since $5 > 4$
- $(1\ 4\ 5\ 2\ 8) \rightarrow (1\ 4\ 2\ 5\ 8)$, Swap since $5 > 2$
- $(1\ 4\ 2\ 5\ 8) \rightarrow (1\ 4\ 2\ 5\ 8)$, since these elements are already in order ($8 > 5$), algorithm does not swap them.

2nd pass:

- $(1\ 4\ 2\ 5\ 8) \rightarrow (1\ 4\ 2\ 5\ 8)$
- $(1\ 4\ 2\ 5\ 8) \rightarrow (1\ 2\ 4\ 5\ 8)$
- $(1\ 2\ 4\ 5\ 8) \rightarrow (1\ 2\ 4\ 5\ 8)$
- $(1\ 2\ 4\ 5\ 8) \rightarrow (1\ 2\ 4\ 5\ 8)$

Will do a third pass since a swap took place in 2nd pass, but the array is sorted!

A lower bound for simple sorting algorithms

- **An inversion** : an ordered pair (A_i, A_j) such that $i < j$ but $A_i > A_j$
- **Example:** 10, 6, 7, 15, 3, 1 Inversions are: (10,6), (10,7), (10,3), (10,1), (6,3), (6,1) (7,3), (7,1) (15,3), (15,1), (3,1)

- Swapping adjacent elements that are out of order removes one inversion.
- A sorted array has no inversions.
- Sorting an array that contains i inversions requires at least i swaps of adjacent elements.

- **Theorem 1:** The average number of inversions in an array of N distinct elements is $N(N - 1)/4$
- **Theorem 2:** Any algorithm that sorts by **swapping adjacent elements** requires $\Omega(N^2)$ time on average.
- For a sorting algorithm to run in **less than quadratic time** it must do something **other than swapping adjacent elements**.


Heapsort

- The priority queue can be used to sort N items by
 - 1 inserting every item into a binary heap OR calling `buildHeap`
 - 2 extracting every item by calling `deleteMin` N times
- An algorithm based on this idea is **heapsort**.
- It is an $O(N \log N)$ worst-case sorting algorithm.

Heapsort


- The main problem with this algorithm is that it uses an extra array for the items exiting the heap.
- We can avoid this problem as follows:
 - After each `deleteMin`, the heap shrinks by 1.
 - Thus the cell that was last in the heap can be used to store the element that was just deleted.
 - Using this strategy, after the last `deleteMin`, the array will contain all elements in **decreasing order**.
- If we want them in increasing order we must use a **max heap**.

Heapsort Example



Max heap after the `buildHeap` phase

Heapsort Example (Cont.)



Heap after the first `deleteMax` operation

- In the implementation of heapsort, an array is utilized as in Binary Heap implementation.
- There are some **minor changes** in the code compared to Binary Heap ADT:
 - Since we use max heap, the comparisons logic is changed from $>$ to $<$.
 - Percolating down needs the current heap size which is lowered by 1 at every deletion.

The Heapsort Sort Algorithm

```
// Standard heapsort.

void heapsort( int [] a )
{
    for( int i = a.length / 2 - 1; i >= 0; i-- ) // buildHeap
        percDown( a, i, a.length );
    for( int j = a.length - 1; j > 0; j-- )
    {
        swapReferences( a, 0, j ); // deleteMax
        percDown( a, 0, j );
    }
}
```

percDown Algorithm

```
// Internal method for heapsort.
// i is the index of an item in the heap.
// Returns the index of the left child.

int leftChild( int i )
{
    return 2 * i + 1;
}

// Internal method for heapsort that is used in
// deleteMax and buildHeap.
// i is the position from which to percolate down.
// n is the logical size of the binary heap.

void percDown( int [] a, int i, int n )
{
    int child;
    int tmp;

    for( tmp = a[ i ] ; leftChild( i ) < n; i = child )
    {
        child = leftChild( i );
        if( child != n-1 && a[ child ] < a[ child+1 ] )
            child++;
        if( tmp < a[ child ] )
            a[ i ] = a[ child ];
        else
            break;
    }
    a[ i ] = tmp;
}
```

Analysis of Heapsort

- It is an $O(N \log N)$ algorithm.
 - First phase: Build heap $O(N)$
 - Second phase: N deleteMax operations: $O(N \log N)$.
- Detailed analysis shows that, the average case for heapsort is poorer than quick sort.
 - Quicksort's worst case however is far worse.
- An average case analysis of heapsort is very complicated, but empirical studies show that there is little difference between the average and worst cases.
 - Heapsort usually takes about twice as long as quicksort.
 - Heapsort therefore should be regarded as something of an insurance policy.
 - On average, it is more costly, but it avoids the possibility of $O(N^2)$.

- Mergesort algorithm is one of the two important divide-and-conquer sorting algorithms (the other one is quicksort).
- It is a recursive algorithm.
 - Divides the list into halves,
 - Sorts each half separately, and
 - Then merges the sorted halves into one sorted array.

Merge Sort Example

theArray:

8	1	4	3	2
---	---	---	---	---

Divide the array in half




Sort the halves

Merge the halves:


- a. $1 < 2$, so move 1 from left half to tempArray
- b. $4 > 2$, so move 2 from right half to tempArray
- c. $4 > 3$, so move 3 from right half to tempArray
- d. Right half is finished, so move rest of left half to tempArray

Temporary array
tempArray:



Copy temporary array back into
original array

theArray:

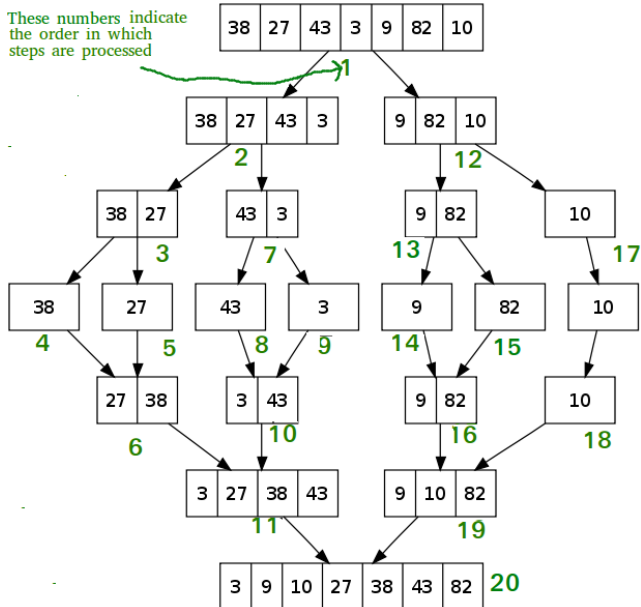


Merge Sort Algorithm

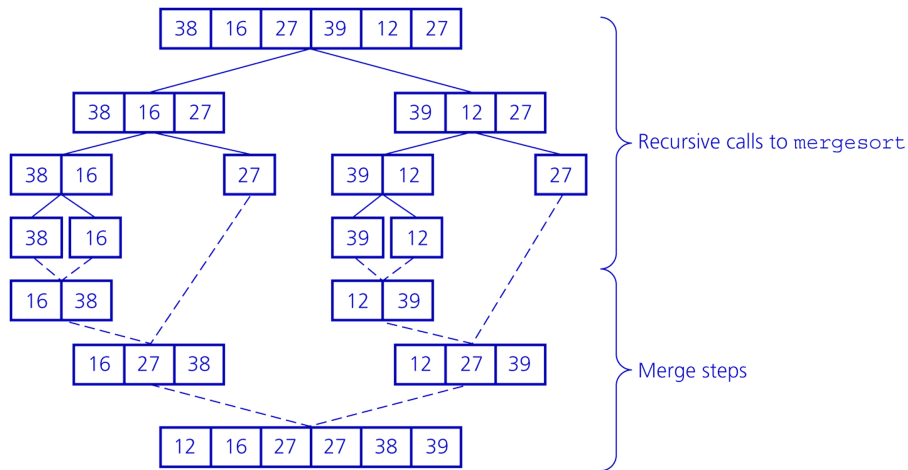
```
mergesort(item_type A[])
{
    n<- length(A), if n<2 return
    middle = n/2
    left: array of size(middle)
    right: array of size (n-middle)
    for i=0:middle-1
        left[i]<-A[i]
    for i=middle:n
        right[i-middle] <- A[i]
    mergesort(left);
    mergesort(right);
    merge(left, right, A);
}

merge(Left, Right, A )
{
    nL <- length(Left)  //nr. of elements on the left
    nR <- length(Right) //nr. of elements on the right
    i,j,k <- 0
    while (i<nL && j<nR) //loop over the left and right
    {
        if(Left[i] < Right[j])
            A[k] <-Left[i], k++, i++ //advance left
        else
            A[k] <-Right[j], k++, j++ //advance right
    }
    while (i<nL) //Copy the remaining
        A[k] <-Left[i], k++, i++
    while (j<nR)
        A[k] <-Right[j], k++, j++
}
```


Merge Sort Example



Merge Sort Example



Mergesort - Analysis of Merge




Merge the halves:

- a. $1 < 4$, so move 1 from theArray[first..mid] to tempArray
- b. $2 < 4$, so move 2 from theArray[first..mid] to tempArray
- c. $8 > 4$, so move 4 from theArray[mid+1..last] to tempArray
- d. $8 > 5$, so move 5 from theArray[mid+1..last] to tempArray
- e. $8 > 6$, so move 6 from theArray[mid+1..last] to tempArray
- f. theArray[mid+1..last] is finished, so move 8 to tempArray

A worst-case instance of the merge step in mergesort

Mergesort - Analysis of Merge (cont.)



Merging two sorted arrays of size k


- Best-case:

- All the elements in the first array are smaller (or larger) than all the elements in the second array.
- The number of moves: $2k + 2k$
- The number of key comparisons: k

Worst-case:

- The number of moves: $2k + 2k$
- The number of key comparisons: $2k - 1$

Mergesort - Analysis



Level 0: mergesort 8 items

Level 1: 2 calls to mergesort with 4 items each

Level 2: 4 calls to mergesort with 2 items each

Level 3: 8 calls to mergesort with 1 item each

Levels of recursive calls to mergesort, given an array of eight items

Mergesort - Analysis

- Assume that N is a power of 2, split into two even halves, for $N=1$ mergesort is constant, 1
- Otherwise, the time to mergesort N numbers is equal to the time to do 2 recursive mergesort of size $N/2$, plus time to merge, which is linear
 $T(1) = 1$, $T(N) = 2T(N/2) + N$
- Let us divide both sides by N
 $T(N)/N = T(N/2)/(N/2) + 1$
- valid for any N that is a power of 2
 $T(N/2)/(N/2) = T(N/4)/(N/4) + 1$ and
 $T(N/4)/(N/4) = T(N/8)/(N/8) + 1$...
 $T(2)/(2) = T(1)/(1) + 1$
- Now add up all the equations. Observe that $T(N/2)/(N/2)$ appear on both sides and cancel. In fact all terms appear on both sides and cancel. This is called *telescoping sum*:
 $T(N)/(N) = T(1)/(1) + \log N$ (there are $\log N$ equations and all 1's add up to $\log N$.
Multiplying by N
 $T(N) = N \log N + N = O(N \log N)$

Mergesort - Analysis

- Mergesort is an extremely efficient algorithm with respect to time.
 - Both worst case and average cases are $O(n \times \log n)$
- Running time depends heavily on relative costs of **comparing** and **moving elements** in the array and the temporary array. These costs are language dependent.
- Mergesort requires an extra array whose size equals to the size of the original array.
- If we use a linked list, we do not need an extra array
 - But, we need space for the links
 - And, it will be difficult to divide the list into half ($O(n)$)

Mergesort for Linked Lists


- Merge sort is often preferred for sorting a linked list. The slow random-access performance of a linked list makes some other algorithms (such as quicksort) perform poorly, and others (such as heapsort) completely impossible.
- MergeSort
 - 1 If head is NULL or there is only one element in the Linked List then return.
 - 2 Else divide the linked list into two halves.
 - 3 Sort the two halves a and b.
 - 4 Merge the two parts of the list into a sorted one.

- Like mergesort, Quicksort is also based on the **divide-and-conquer** paradigm.
- But it uses this technique in a somewhat opposite manner, as all the hard work is done before the recursive calls.
- Basic idea:
 - 1 First, arbitrarily choose an item and partition the array into three groups: those smaller than the chosen item, those equal to the chosen item and those larger than the chosen item,
 - 2 Then, sort the first and last groups recursively,
 - 3 Finally, concatenate three groups.

Algorithm 1 Quicksort

- 1: Let S be the input set.
 - 2: **if** $|S| = 0$ or $|S| = 1$ **then**
 return
 - 3: Pick an element v in S . Call v the **pivot**.
 - 4: Partition $S - \{v\}$ into two disjoint groups:
 $S_1 = \{x \in S - \{v\} \mid x < v\}$
 $S_2 = \{x \in S - \{v\} \mid x \geq v\}$
 return { quicksort(S_1), v , quicksort(S_2) }
-

Quicksort Illustrated




Issues To Consider

- How to pick the pivot?
 - Many methods (discussed later)
- How to partition?
 - Several methods exist.
 - The one we consider is known to give good results and to be easy and efficient.

We discuss the partition strategy first.

Partitioning Strategy


- For now, assume that $\text{pivot} = A[(\text{left} + \text{right})/2]$.
- We want to partition array $A[\text{left} .. \text{right}]$.
- First, get the pivot element out of the way by swapping it with the last element (swap pivot and $A[\text{right}]$).
- Let i start at the first element and j start at the next-to-last element ($i = \text{left}$, $j = \text{right} - 1$)



Partitioning Strategy (Cont.)

- Want to have

- $A[k] \leq \text{pivot}$, for $k < i$
- $A[k] \geq \text{pivot}$, for $k > j$




- When $i < j$

- Move i right, skipping over elements smaller than the pivot
- Move j left, skipping over elements greater than the pivot
- When both i and j have stopped
 - $A[i] \geq \text{pivot}$ (i is pointing to a larger element)
 - $A[j] \leq \text{pivot}$ (j is pointing to a smaller element)
 - $\Rightarrow A[i]$ and $A[j]$ should now be swapped



Partitioning Strategy (Cont.)

- When i and j have stopped and i is to the left of j (thus legal)
 - Swap $A[i]$ and $A[j]$
 - The large element is pushed to the right and the small element is pushed to the left
 - After swapping
 - $A[i] \leq \text{pivot}$
 - $A[j] \geq \text{pivot}$
 - Repeat the process until i and j cross




Partitioning Strategy (Cont.)

- When i and j have crossed
 - swap $A[i]$ and pivot
- Result:
 - $A[k] \leq \text{pivot}$, for $k < i$
 - $A[k] \geq \text{pivot}$, for $k > i$



swap $A[i]$ and pivot



Break!

j i

Pivot Strategies

- **First element:**

- Bad choice if input is sorted or in reverse sorted order
- Bad choice if input is nearly sorted

- **Random:**

- Perfectly safe, unless the random number generator has a flaw
- Random number generation may be costly

- **Median-of-three:**

- The best choice would be the median of the array but not feasible.
- A good estimate is to choose the median of the left, right, and center elements

Quicksort algorithm

It will be implemented in PS

Analysis of Quicksort

- **Worst case:** pivot is the smallest (or largest) element all the time.

$$T(N) = T(N - 1) + cN$$

$$T(N - 1) = T(N - 2) + c(N - 1)$$

$$T(N - 2) = T(N - 3) + c(N - 2)$$

...

$$T(2) = T(1) + c(2)$$

$$T(N) = T(1) + c \sum_{i=2}^N i \rightarrow O(N^2)$$

- **Best case:** pivot is the median

$$T(N) = 2T(N/2) + cN$$

$$T(N) = cN \log N + N \rightarrow O(N \log N)$$

Quicksort: Average case

- Assume each of the sizes for S_1 are equally likely.
- $0 \leq |S_1| \leq N - 1$.

$$T(N) = \left(\frac{1}{N} \sum_{i=0}^{N-1} [T(i) + T(N - i - 1)] \right) + cN$$

$$T(N) = \left(\frac{2}{N} \sum_{i=0}^{N-1} T(i) \right) + cN$$

$$NT(N) = \left(2 \sum_{i=0}^{N-1} T(i) \right) + cN^2$$

$$(N - 1)T(N - 1) = \left(2 \sum_{i=0}^{N-2} T(i) \right) + c(N - 1)^2$$

$$NT(N) - (N - 1)T(N - 1) = 2T(N - 1) + 2cN - c$$

$$NT(N) = (N + 1)T(N - 1) + 2cN$$

Quicksort: Average case (Cont.)

$$NT(N) = (N + 1)T(N - 1) + 2cN$$

Divide equation by $N(N + 1)$

$$\frac{T(N)}{N+1} = \frac{T(N-1)}{N} + \frac{2c}{N+1}$$

Telescope;

$$\begin{aligned}\frac{T(N-1)}{N} &= \frac{T(N-2)}{N-1} + \frac{2c}{N} \\ \frac{T(N-2)}{N-1} &= \frac{T(N-3)}{N-2} + \frac{2c}{N-1}\end{aligned}$$

$$\dots$$
$$\frac{T(2)}{3} = \frac{T(1)}{2} + \frac{2c}{3}$$

$$\frac{T(N)}{N+1} = \frac{T(1)}{2} + 2c \sum_{i=3}^{N+1} \frac{1}{i}$$

Quicksort: Average case (Cont.)

$$2c \sum_{i=3}^{N+1} \frac{1}{i} = 2c(H_{N+1} - \frac{3}{2})$$

$$T(N) = (N+1)(\frac{T(1)}{2} + 2c(H_{N+1} - \frac{3}{2}))$$

$H_N \approx \log_e(N) + \gamma + \frac{1}{2N}$ ($\gamma = 0.577215664901$ (Euler-Mascheroni Constant))

$$T(N) \approx (N+1) \left[\frac{T(1)}{2} + 2c \left((\log_e(N+1) + \gamma + \frac{1}{2(N+1)}) - \frac{3}{2} \right) \right]$$

$$T(N) \rightarrow O(N \log N)$$

How fast can we sort?

- Heapsort, Mergesort, and Quicksort all run in $O(N \log N)$ best case running time.
- Can we do any better?


The Answer is No! (if using comparisons only)

- Our basic assumption: we can only compare two elements at a time – how does this limit the run time?
- Suppose you are given N elements
 - Assume no duplicates – any sorting algorithm must also work for this case
- How many possible orderings can you get?

How many possible orderings?

- Example: a, b, c ($N = 3$)
- Orderings:
 - ① a b c
 - ② b c a
 - ③ c a b
 - ④ a c b
 - ⑤ b a c
 - ⑥ c b a
- 6 orderings = $3 \times 2 \times 1 = 3!$
- For N elements: $N!$ orderings

A Decision Tree



Leaves contain possible orderings of a , b , c

Decision Trees and Sorting

- A Decision Tree is a Binary Tree such that:
 - Each node = a set of orderings
 - Each edge = 1 comparison
 - Each leaf = 1 unique ordering
 - How many leaves for N distinct elements?
- Every algorithm that sorts by using only comparisons can be represented by a decision tree
 - Only 1 leaf has sorted ordering
 - Finds correct leaf by following edges (= comparisons)
- Run time \geq maximum number of comparisons
 - Depends on: depth of decision tree
 - What is the depth of a decision tree for N distinct elements?

Lower Bound on Comparison-Based Sorting

- Suppose you have a binary tree of depth d . How many leaves can the tree have?
 - e.g. depth $d = 1 \rightarrow$ at most 2 leaves,
 - $d = 2 \rightarrow$ at most 4 leaves, etc.
- A binary tree of depth d has at most 2^d leaves
- Number of leaves $L \leq 2^d \rightarrow d \geq \log L$
- Decision tree has $L = N!$ leaves \rightarrow its depth $d \geq \log(N!)$

Lower Bound on Comparison-Based Sorting

- Stirling's approximation: $N! \approx \sqrt{2\pi N} \left(\frac{N}{e}\right)^N$
- $\log(N!) \approx \log\left(\sqrt{2\pi N} \left(\frac{N}{e}\right)^N\right)$
$$= \log\left(\sqrt{2\pi N}\right) + \log\left(\left(\frac{N}{e}\right)^N\right)$$
$$= \frac{1}{2}\log(2\pi N) + N(\log(N) - 1) \rightarrow \Omega(N \log N)$$
- **Conclusion:** Any sorting algorithm **based on comparisons between elements** requires $\Omega(N \log N)$ comparisons

Comparison of Sorting Algorithms

Algorithm	Worst case	Average case
Selection sort	$O(N^2)$	$O(N^2)$
Bubble sort	$O(N^2)$	$O(N^2)$
Insertion sort	$O(N^2)$	$O(N^2)$
Heapsort	$O(N \log N)$	$O(N \log N)$
Mergesort	$O(N \log N)$	$O(N \log N)$
Quicksort	$O(N^2)$	$O(N \log N)$

Sorting in linear time

- Comparison sort:
 - Lower bound: $\Omega(n \log n)$.
- Non comparison sort:
 - Bucket sort, radix sort
 - They can sort in linear time (under certain assumptions).

Bucket Sort

- Assumption: The input A_1, A_2, \dots, A_N consists of only positive integers smaller than M (obviously extensions are possible)
- Algorithm:
 - Keep an array called *count* (of size M), which is initialized to all 0s. (*count* has M buckets which are all empty)
 - When A_i is read, increment $\text{count}[A_i]$ by 1.
 - After all the input is read, scan the *count* array, printing out the representation of the sorted array.
- Running time will be $O(M + N)$.
If M is $O(N)$ then running time will be $O(N)$.
- also called counting sort

- Assumption: Input of size n is uniformly distributed over an interval.
- Above algorithm can be extended as follows:
 - Divide input range into equal-sized subintervals (buckets).
 - Distribute n numbers into buckets.
 - Sort items in each bucket with a generic sorting algorithm.
 - Go through buckets in order to create sorted array.
- Worst case running time will be $O(N^2)$ but if number of buckets is $O(N)$ and the numbers are not distributed non-uniformly to the buckets, then average running time will be $O(N)$.

Radix Sort

- Origin: Herman Hollerith's card-sorting machine for the 1890 U.S. Census
- Digit-by-digit sort.
- Hollerith's original (not optimal) idea: sort on most-significant digit first.
- Good idea: Sort on least-significant digit first with auxiliary stable sort.
- **Stable Sort Property:** The relative order of any two items with the same key is preserved after the execution of the algorithm.

Radix Sort Algorithm

Algorithm 2 RadixSort(A, d)

1: **for** $i \leftarrow 1$ **to** d **do**

 use stable BucketSort to sort array A on digit i .

- **Lemma:** Given n d -digit numbers in which each digit can take on up to k possible values, RadixSort correctly sorts these numbers in $\Theta(d(n + k))$ time.
 - If d is constant and $k = O(n)$, then time is $\Theta(n)$.

Radix Sort Example

INITIAL ITEMS:	064, 008, 216, 512, 027, 729, 000, 001, 343, 125
SORTED BY 1's digit:	000, 001, 512, 343, 064, 125, 216, 027, 008, 729
SORTED BY 10's digit:	000, 001, 008, 512, 216, 125, 027, 729, 343, 064
SORTED BY 100's digit:	000, 001, 008, 027, 064, 125, 216, 343, 512, 729